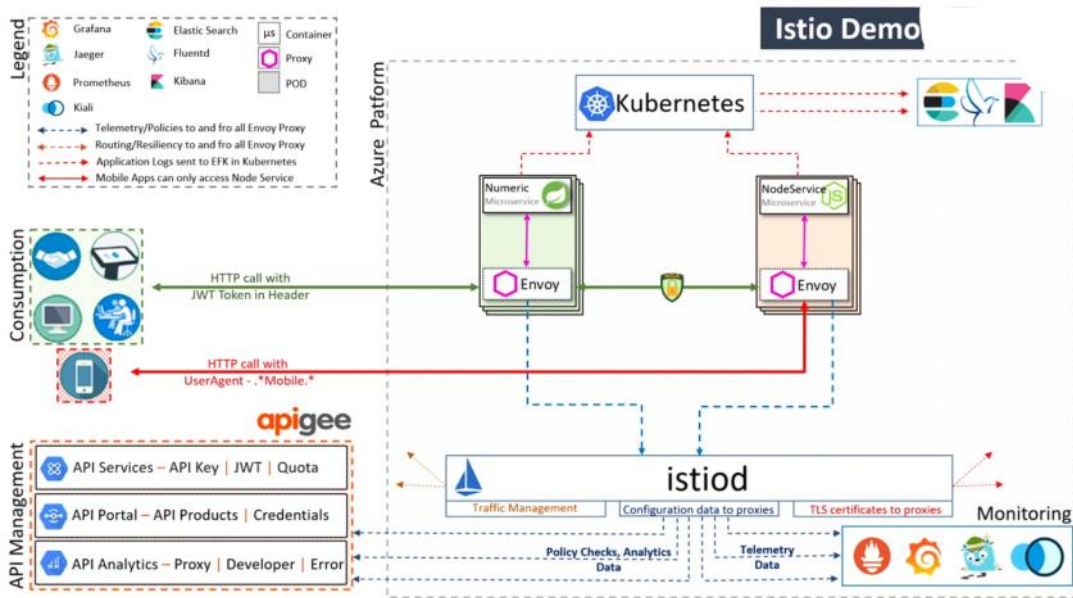


# Istio service mesh - Production deployment

## Istio architecture



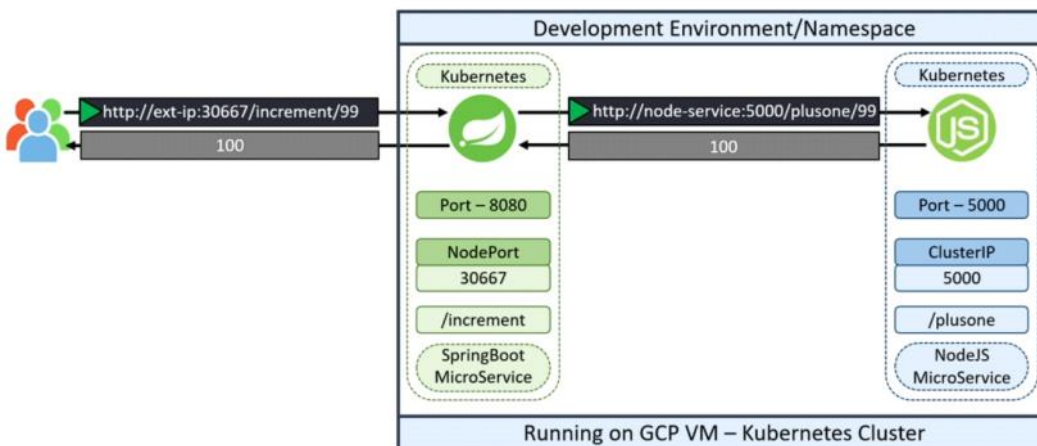
### Overview:

Istio is an open-source service mesh platform that provides a unified control plane for managing microservices. Istio enables developers to add advanced traffic management, security, and observability to their microservices without requiring changes to the application code.

### Reference:

[Giới thiệu Istio - Istio là gì \(devopsvn.tech\)](https://devopsvn.tech/gioi-thieu-istio-istio-la-gi/)

### Application architecture:



## Istio Installation

```
curl -Ls https://istio.io/downloadIstio | ISTIO_VERSION=1.17.2 sh -  
cd istio-1.17.2  
export PATH=$PWD/bin:$PATH  
istioctl install --set profile=demo -y && kubectl apply -f samples/addons
```

Check all node,pod,svc in istio namespace:

```
kubectl -n istio-system get all
```

You can access services inside istio using a web browser by modifying the svc in istio from clusterIP to nodeport with command below.

```
kubectl -n istio-system edit svc kiali
```

However, this approach is not recommended in a practical environment.

Another way is using port-forwarding. You can create script like below.

port-forwarding.sh

```
kubectl -n istio-system port-forward deploy/kiali 20001:20001 --address 192.168.207.129 &
kubectl -n istio-system port-forward deploy/grafana 3000:3000 --address 192.168.207.129 &
kubectl -n istio-system port-forward deploy/prometheus 9090:9090 --address 192.168.207.129 &
kubectl -n istio-system port-forward deploy/jaeger 16686:16686 --address 192.168.207.129 &
```

Now, we create bash shell file to setup production environment

```
#!/bin/bash
# Set the namespace name
namespace="prod"
# Check if the namespace exists
if kubectl get namespace "${namespace}" > /dev/null 2>&1; then
    # Check if the namespace has the istio-injection label
    if ! kubectl get namespace "${namespace}" -o jsonpath='{.metadata.labels.istio-injection}' | grep -q enabled; then
        # Add the istio-injection label to the namespace
        kubectl label namespace "${namespace}" istio-injection=enabled
        # Restart all deployments in the namespace to apply the label
        kubectl rollout restart deployment -n "${namespace}"
    else
        echo "namespace ${namespace} already has istio-injection label"
    fi
else
    # Create the namespace and add the istio-injection label
    kubectl create namespace "${namespace}"
    kubectl label namespace "${namespace}" istio-injection=enabled
fi
# Get the name and tag of the Docker image
image_info=$(docker images --format '{{.Repository}}:{{.Tag}}' | grep -i node-service)
if [ -n "${image_info}" ]; then
    image_name=$(echo "${image_info}" | cut -d ':' -f 1)
    image_tag=$(echo "${image_info}" | cut -d ':' -f 2)
    # Check if the deployment already exists
    if ! kubectl get deployment node-app -n "${namespace}" > /dev/null 2>&1; then
        # Create the deployment with 2 replicas and the specified image
        kubectl create deployment node-app --image "${image_name}:${image_tag}" --replicas 2 -n "${namespace}"
    else
        echo "deployment node-app already exists"
    fi
    # Check if the service already exists
    if ! kubectl get service node-service -n "${namespace}" > /dev/null 2>&1; then
        # Create a service for the deployment on port 5000
        kubectl expose deployment node-app --name node-service --port 5000 -n "${namespace}"
    else
        echo "service node-service already exists"
    fi
else
    echo "Docker image for node-service does not exist"
fi
```

The source code below is a shell script that performs the following tasks:

- It checks if the namespace prod exists in the Kubernetes cluster using the kubectl command. A namespace is a way to group and isolate resources within a cluster.
- If the namespace prod exists, it checks if the label istio-injection=enabled exists in the namespace. A label is a key-value pair that can be attached to any resource for identification and selection. Istio is a service mesh that provides traffic management, security and observability

for microservices. The label `istio-injection=enabled` enables automatic sidecar injection for the pods in the namespace. A sidecar is a container that runs alongside the main application container and provides additional functionality such as proxying, logging and monitoring.

- If the label `istio-injection=enabled` does not exist in the namespace `prod`, it adds the label using the `kubectl` command.
- If the namespace `prod` does not exist, it creates the namespace using the `kubectl` command and adds the label `istio-injection=enabled` to it.
- It gets the name of the Docker image `node-service` using the `docker` command. A Docker image is a packaged version of an application and its dependencies that can run on any platform that supports Docker. Node-service is an example of a Node.js application that serves HTTP requests on port 5000.
- It checks if the Docker image `node-service` exists by checking if the image name is not empty.
- If the image exists, it gets the image tag using the `docker` command. A tag is a way to identify a specific version of an image.
- It creates a deployment and a service for the image using the `kubectl` command. A deployment is a resource that manages a set of pods that run the same image and can be scaled up or down. A service is a resource that exposes a set of pods to other pods or external clients using a stable IP address and port number.
- It specifies the namespace `prod`, the image name and tag, and the number of replicas (2) for the deployment. It also specifies the name (`node-service`), the port (5000) and the type (ClusterIP) for the service. ClusterIP means that the service is only accessible within the cluster.
- If the image does not exist, it prints an message.

You can manual check and apply `istio-injection` following below command:

```
sudo kubectl get all -n prod
sudo kubectl label ns prod istio-injection=enabled
sudo kubectl get ns --show-labels
kubectl rollout restart deployment node-app -n prod
```

Next. We create file `k8s_PROD-deployment_service.yaml` to deploy application to production enviroment.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app: devsecops
    name: devsecops
spec:
  replicas: 3
  selector:
    matchLabels:
      app: devsecops
  strategy: {}
  template:
    metadata:
      labels:
        app: devsecops
    spec:
      serviceAccountName: default
      volumes:
        - name: vol
          emptyDir: {}
      containers:
        - image: replace
          name: devsecops-container
          volumeMounts:
            - mountPath: /tmp
              name: vol
      securityContext:
        capabilities:
          drop:
            - NET_RAW
        runAsUser: 100
        runAsNonRoot: true
        readOnlyRootFilesystem: true
        allowPrivilegeEscalation: false
      resources:
        requests:
          memory: "256Mi"
          cpu: "200m"
```

```

    limits:
      memory: "512Mi"
      cpu: "500m"
---
apiVersion: v1
kind: Service
metadata:
  labels:
    app: devsecops
  name: devsecops-svc
spec:
  ports:
    - port: 8080
      protocol: TCP
      targetPort: 8080
  selector:
    app: devsecops
  type: ClusterIP

```

The resources section in the YAML file specifies the resource requirements and limits for the container. In this case, the container has a request of 256Mi memory and 200m CPU. This means that Kubernetes will try to schedule the container on a node that has at least 256Mi memory and 200m CPU available. If there are no nodes with enough resources available, the container will remain in a pending state.

The container also has a limit of 512Mi memory and 500m CPU. This means that the container will not be allowed to use more than 512Mi memory and 500m CPU even if more resources are available on the node.

Setting requests and limits is important for a number of reasons. First, it helps to ensure that your containers don't over-consume resources and cause other containers to fail. Second, it can help to improve the performance of your applications by ensuring that they have the resources they need to run efficiently. Third, it can help to prevent your applications from being killed by Kubernetes if they try to use too many resources.

A securityContext that takes away the NET\_RAW capability from our application can help us avoid dns spoofing.

Reference: <https://blog.aquasec.com/dns-spoofing-kubernetes-clusters>

Next, Create file k8s-PROD-deployment-rollout-status.sh like

```

#!/bin/bash
sleep 60s
if [[ $(kubectl -n prod rollout status deploy ${deploymentName} --timeout 5s) != *"successfully rolled out"* ]];
then
  echo "Deployment ${deploymentName} Rollout has Failed"
  kubectl -n prod rollout undo deploy ${deploymentName}
  exit 1;
else
  echo "Deployment ${deploymentName} Rollout is Success"
fi

```

Add k8s deployment production stage into Jenkins file

```

stage('K8S Deployment - PROD') {
  steps {
    parallel(
      "Deployment": {
        withKubeConfig([credentialsId: 'kubeconfig']) {
          sh "sed -i 's#replace#${imageName}#g' k8s_PROD-deployment_service.yaml"
          sh "kubectl -n prod apply -f k8s_PROD-deployment_service.yaml"
        }
      },
      "Rollout Status": {
        withKubeConfig([credentialsId: 'kubeconfig']) {
          sh "bash k8s-PROD-deployment-rollout-status.sh"
        }
      }
    )
  }
}

```

```
}  
}
```

Run build. After build successful you can check connection like below.

## Istio PeerAuthentication

Next, we create PeerAuthentication istio-system with yaml file below.

peer-auth.yaml

```
apiVersion: security.istio.io/v1beta1  
kind: PeerAuthentication  
metadata:  
  name: default  
  namespace: istio-system  
spec:  
  mtls:  
    mode: DISABLE
```

This YAML file is a configuration file for Istio's PeerAuthentication resource, which is used to configure mutual TLS (mTLS) authentication between services in an Istio service mesh. The file specifies the following:

- apiVersion: The version of the Istio security API being used.
- kind: The type of Istio resource being defined, which is PeerAuthentication in this case.
- metadata: The metadata for the PeerAuthentication resource, which includes the name of the resource (default) and the namespace in which it is defined (istio-system).
- spec: The specification for the PeerAuthentication resource, which includes the mtls field that specifies the mTLS mode.

The mtls field specifies the mTLS mode for the PeerAuthentication resource, which is set to DISABLE in this case. This means that mTLS authentication is disabled for all services in the Istio service mesh.

In general, there are three mTLS modes that can be set for a PeerAuthentication resource:

- STRICT: All requests between services must use mTLS authentication.
- PERMISSIVE: Requests between services can use either mTLS authentication or plaintext (non-encrypted) communication.
- DISABLE: mTLS authentication is disabled for all services in the Istio service mesh.

The PeerAuthentication resource can be used to configure mTLS authentication for specific services or for all services in the Istio service mesh.

Run kubectl apply:

```
kubectl apply -f peer-auth.yaml -n istio-system
```

Check mtls mode:

```
sudo kubectl get pa -n istio-system  
NAME    MODE    AGE  
default DISABLE  9s
```

You can change mtls mode with command below:

```
kubectl edit pa -n istio-system  
spec:  
  mtls:  
    mode: STRICT # write mode you want
```

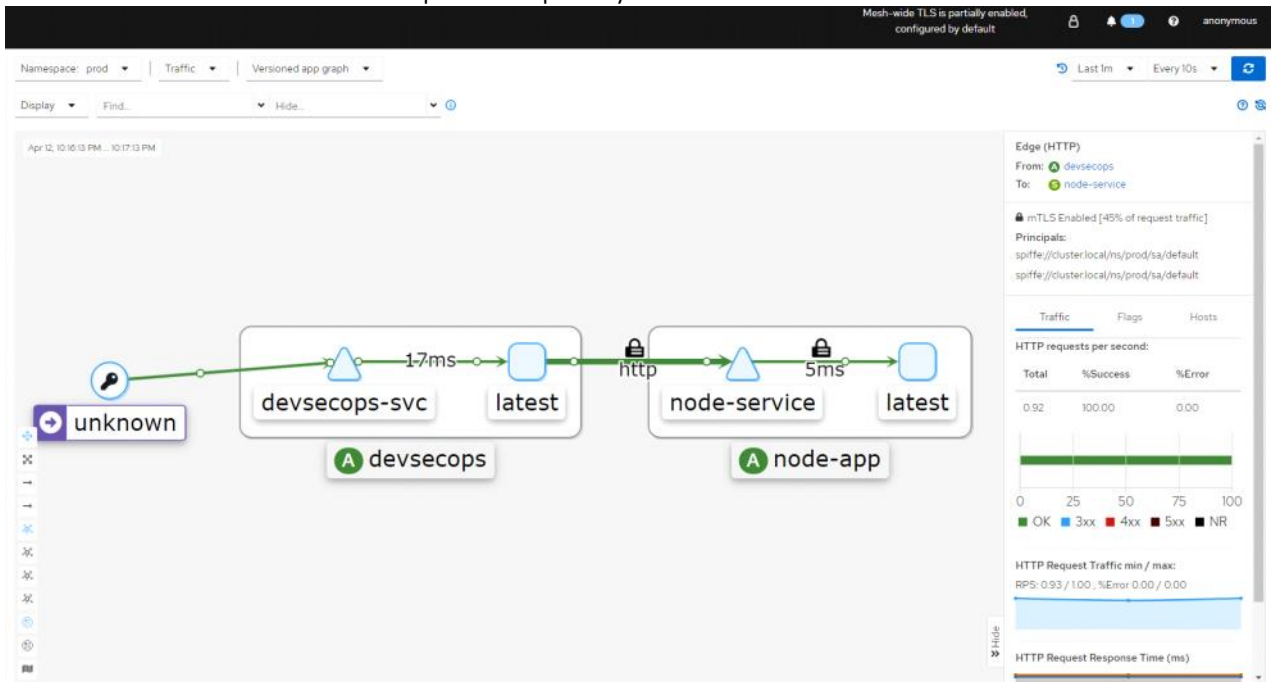
Reference: [https://istio.io/latest/docs/reference/config/security/peer\\_authentication/](https://istio.io/latest/docs/reference/config/security/peer_authentication/)

We will run command below to test connection

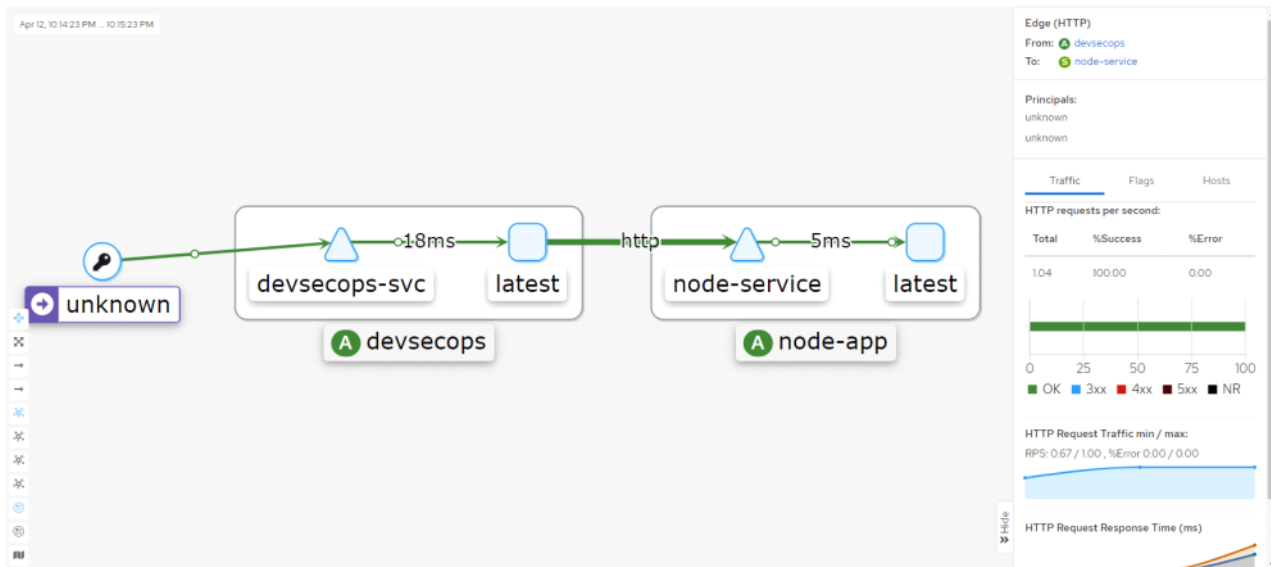
```
sudo kubectl -n prod get svc  
NAME      TYPE      CLUSTER-IP    EXTERNAL-IP  PORT(S)  AGE  
devsecops-svc ClusterIP  10.103.98.41  <none>       8080/TCP  148m  
node-service ClusterIP  10.110.155.212 <none>       5000/TCP  148m
```

```
while true; do curl -s 10.103.98.41:8080/increment/99; echo; sleep 1; done
```

You can go to kiali dashboard, select graph to know how connection inside out application work.  
PERMISSIVE MODE -we can see at the top there is a partially enabled TLS text

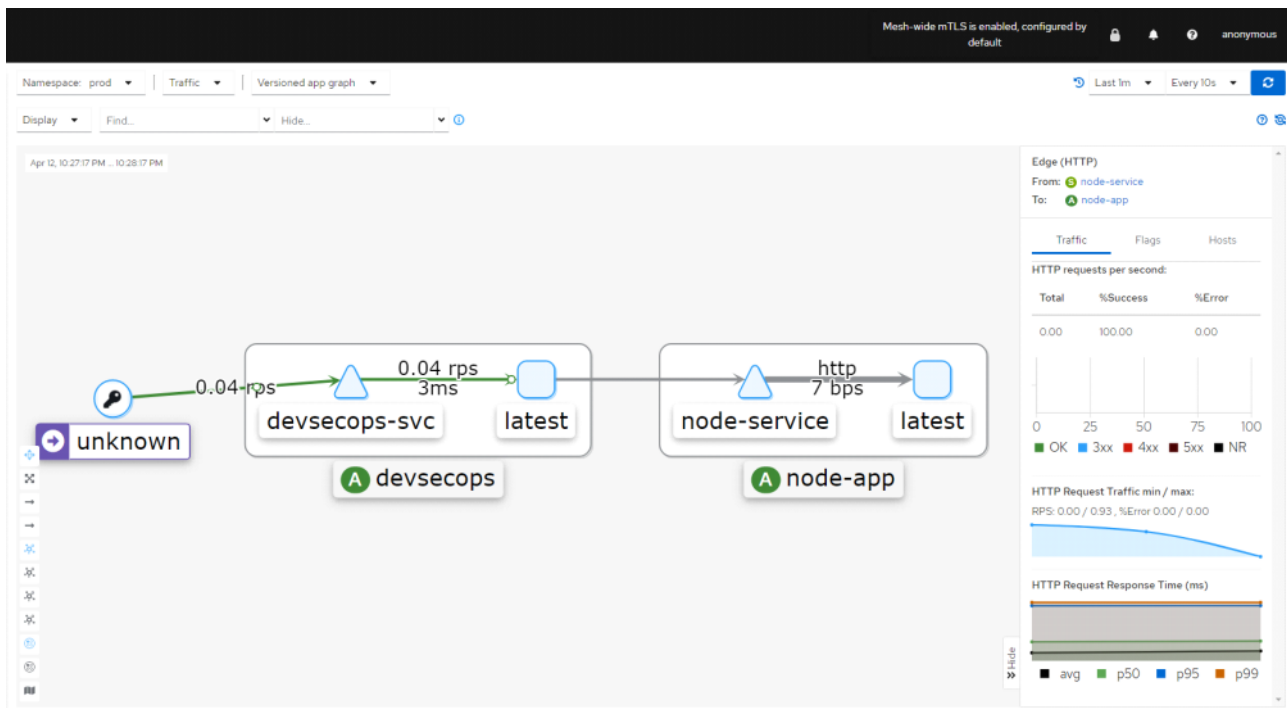


DISABLE MODE -when we change the mode to disabled, traffic including mTLS will drop until all traffic is no longer encrypted

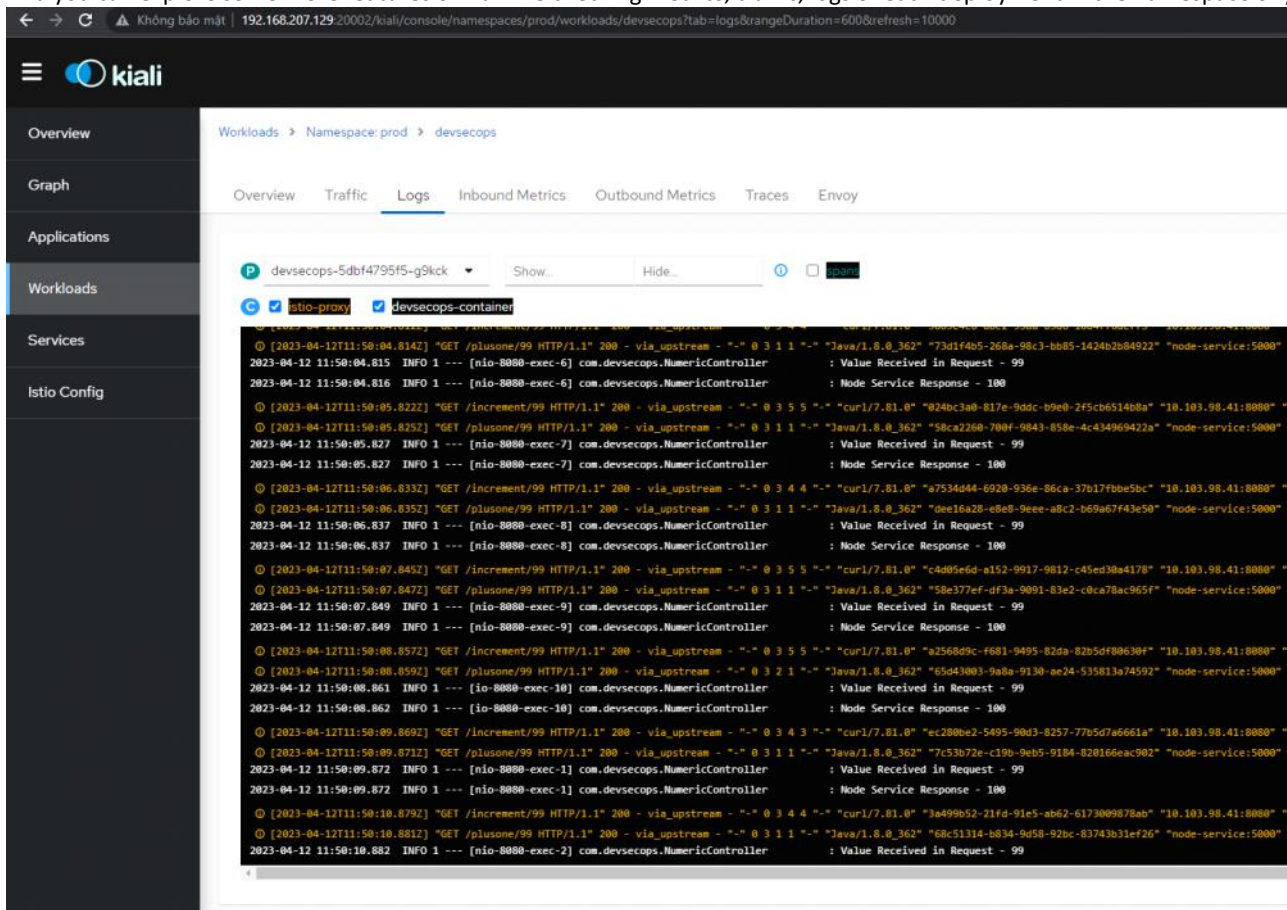


STRICT MODE: all traffic is encrypted, authenticated and does not allow outside traffic to go inside





And you can explore some more features of kiali like checking metrics, traffic, logs of each deployment in the namespace of your choice.



## ISTIO INGRESS

In Istio, an Ingress is a component that enables external traffic to access services within the mesh. Istio Ingress provides traffic management

features such as load balancing, TLS termination, and routing. It allows you to define rules for routing traffic to specific services based on various criteria such as URL path, headers, and source IP address.

A VirtualService, on the other hand, is a Kubernetes custom resource that is used to define rules for how traffic should be routed to a specific service within the Istio service mesh. It provides more advanced traffic management features than Ingress, such as traffic splitting, fault injection, and retries. With VirtualServices, you can define rules to send traffic to different versions of a service or to specific instances of a service based on various criteria such as HTTP headers or source IP addresses.

In summary, Istio Ingress provides a way for external traffic to enter the service mesh and provides basic routing capabilities, while VirtualService provides more advanced traffic management features for routing traffic within the mesh.

Reference: <https://istio.io/latest/docs/tasks/traffic-management/ingress/ingress-control/>  
<https://istio.io/latest/docs/reference/config/networking/virtual-service/>

Now we go to create istio ingress and virtualService according to the yaml file below.

```
apiVersion: networking.istio.io/v1alpha3 # The API version for Istio resources
kind: Gateway # The kind of resource
metadata: # The metadata for the resource
  name: devsecopsgateway # The name of the Gateway
  namespace: prod # The namespace where the Gateway is deployed
spec: # The specification for the Gateway
  selector: # The selector for the Gateway
    istio: ingressgateway # The label of the Istio ingress gateway pod
  servers: # The list of servers for the Gateway
    - port: # The port configuration for the server
        number: 80 # The port number
        name: http # The port name
        protocol: HTTP # The port protocol
      hosts: # The list of hosts for the server
        - "*" # A wildcard host that matches any domain
--- # The separator for the next document
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: devsecopsnumeric # The name of the VirtualService
  namespace: prod
spec: # The specification for the VirtualService
  hosts:
    - "*"
  gateways: # The list of gateways for the VirtualService
    - devsecopsgateway # The name of the Gateway from the first document
  http: # The list of HTTP routes for the VirtualService
    - match: # The list of match conditions for the route
        - uri: # The URI match condition
            prefix: /increment # The URI prefix to match
        - uri:
            exact: / # The exact URI to match
      route: # The list of destinations for the route
        - destination: # The destination configuration
            host: devsecopssvc # The host name of the destination service
            port:
              number: 8080
```

This is a YAML file that contains two Kubernetes resources: a Gateway and a VirtualService.

The Gateway resource is named devsecops-gateway and is of type networking.istio.io/v1alpha3. It has a selector that specifies the Istio ingress gateway as the default controller. It also has one server that listens on port 80 and accepts HTTP traffic from any host.

The VirtualService resource is named devsecops-numeric and is also of type networking.istio.io/v1alpha3. It specifies that all hosts should be matched and that the gateway to use is devsecops-gateway. It also specifies an HTTP route that matches URIs with prefix /increment or exact URI / and sends traffic to a destination service named devsecops-svc listening on port 8080.

URI prefix and exact URI are two different concepts in web development. A URI is a string of characters that identifies a name or a resource



on the internet. A URI prefix is a matching method based on the prefix of the URL path separated by the slash character. The matching is case-sensitive and performed on each segment of the path. On the other hand, an exact URI is a matching method that matches the entire URL path, including the query parameters and the fragment identifier. In web development, these concepts are used in routing and URL mapping to direct incoming requests to the appropriate controller or handler. The difference between URI prefix and exact URI is that URI prefix matches the URL path based on the prefix of the path, while exact URI matches the entire URL path, including the query parameters and the fragment identifier.

```
kubectl get vs,gateway -n prod
NAME                                GATEWAYS          HOSTS  AGE
virtualservice.networking.istio.io/devsecops-numeric ["devsecops-gateway"] ["*"] 6d23h
NAME                                AGE
gateway.networking.istio.io/devsecops-gateway 6d23h

kubectl -n istio-system get svc istio-ingressgateway
NAME            TYPE           CLUSTER-IP    EXTERNAL-IP  PORT(S)
istio-ingressgateway LoadBalancer  10.99.213.128 <pending>
15021:30561/TCP,80:30886/TCP,443:32190/TCP,31400:32600/TCP,15443:30266/TCP
```

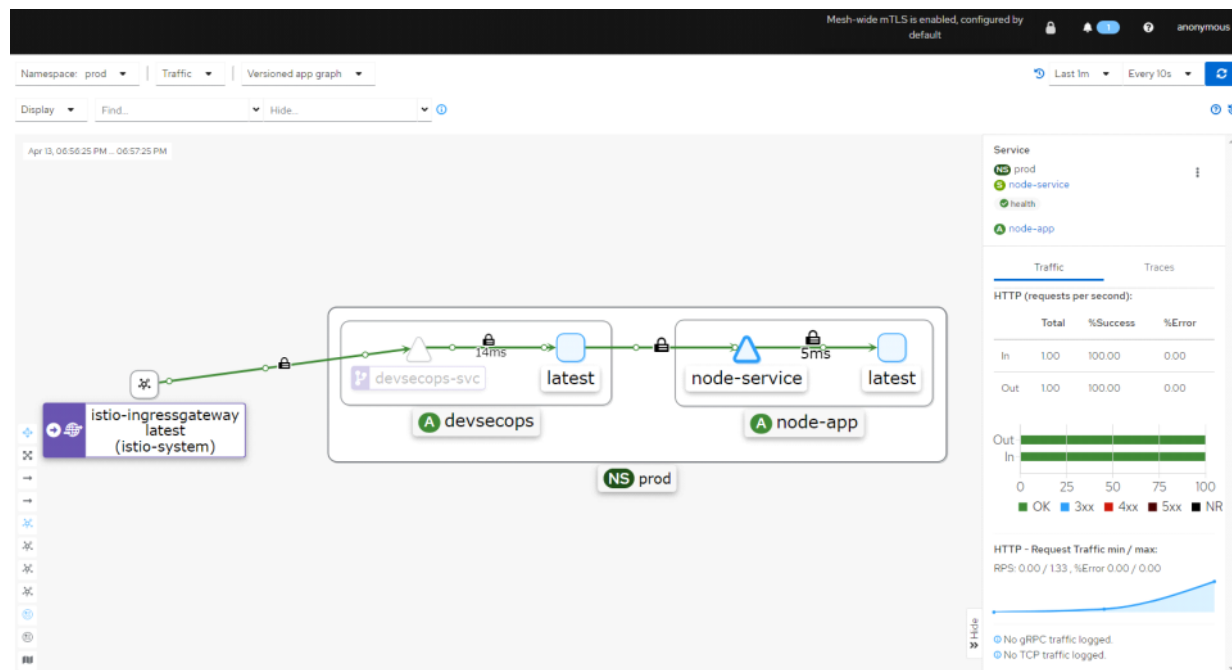
Now we try curl to test connection

```
curl localhost:30886/increment/22
23
```

```
curl localhost:30886/
Kubernetes DevSecOps
```

We can use bash below and open kiali to see connection inside

```
while true; do curl -s localhost:30886/increment/22; echo; sleep 1; done
```



If we remove url extract in virtualService

```
- uri:
  exact: /
```

By using kubectl edit vs -n prod In file and run curl localhost:30886/ to test connection again, we cannot receive response.